

Microsoft DevOps Solutions: Planning Deployment Environment Strategies

UNDERSTANDING RELEASE STRATEGIES



Chris B. Behrens

SOFTWARE ARCHITECT

@chrisbbehrens



A Different Kind of Course

Narrower focus and shallower depth

Deployment is my jam – check out my other courses



Release Strategies on the Test



A finite number of
patterns



The exam's coverage
is good for the real
world



These are patterns
you'll need in real life

The Pattern List



Blue – Green Deployments



Rolling Deployments



Ring or Canary



Blue-Green Deployment



Two separate production-possible environments

“Production” is a designation, not a dedicated set of resources

The designation is a network designation

You have a Staging/Test/Verification environment

You’re almost there already

Without Blue-Green, you push Staging code to production, but with production transforms



The Other Fifth of the Way to Blue-Green

Flip your network
to point to
Staging

What WAS
production
becomes staging,
and vice versa

This buys you
rollback...
Sort of



Network Considerations for Deployment



“simply flip your network”



**It's 10 AM in the morning somewhere
for your audience**



Draining

Gets your users from the old server to the new one

Azure Application Gateway

The old server stops getting new connections

This is a widely used pattern



We want a hot backup in case our deployment didn't go as planned, and we want to make rollback as simple as possible



Rolling Deployments



More than one server



Deployment on a gradual basis



Backward Compatibility

If you think about
it for a minute

Backward
compatible with the
previous version of
the code


Otherwise, the
mixed environment
will break

To remove a
function, remove
the possibility of
request first


And then remove
the function on the
next deployment



The Key Concept in Rolling Deployments



Risk Mitigation for Rolling Deployments



Introduces risk
gradually



So only 10% of your
servers can be in
trouble



You obviously need to
have ten servers to be
able to deploy to 10%
of them




Canary / Ring Deployments

**A more sophisticated version
of rolling deployments**

**The difference in versions is
the point**



What's the Canary?




Canaries were kept in mines

Canaries' systems were very sensitive to CO and other toxic gases

When the canaries fell over, it was time to get everybody out of the mine

Miners took good care of their canaries, and so should you


Who's the Canary?



Risk Mitigation for Canary Deployments



Find yourself some trusted users



And find out what they think of your
new release

(Not Much) Azure Tooling for All of This



Draining is supported in the interface



There's no "Roll Deployment" button



You'll mostly have to roll Canary deployments yourself



With one minor exception



Blue Green Tooling with Deployment Slots

**Deployment slots
enable the Blue-
Green pattern**

**But they can only
be a part of the
solution**

**Taking into
account the rest
of your
infrastructure**



The Essential Problem of Database Availability



1. Reconcile the existing state with the changes
2. Blow away what's there and replace it with the changes
3. Create a brand-new environment and replace the old one

The Problem with Databases



Databases cannot be recreated from nothing (ex nihilo) every time

To do that we would need:

- The schema
- The data


Does your data belong in version control? Maybe not

Getting sensitive stuff out of version control is a huge pain

And if your database isn't read-only, this is all academic



What's Happening Here



Read-only databases are rarely useful

Your customers are constantly rewriting the code necessary to regenerate your database

Database availability is the fixed point you've got to work around

Backward compatibility helps, but...

It doesn't maintain rollback-ability

Some schema changes are truly irreversible, like a column drop



Solutions to This

**“No irreversible
schema changes”**

**Better get a good
layer of abstraction
over that**

**Transitional
deployment to
bridge the gap**

**Code which can
run on EITHER
version**

Key to the version



Schema-on-Write



Transactional database are schema-on-write



Schema enforced at INSERT or UPDATE




This works well for transactional considerations




But it binds a version of the code tightly to the schema



Schema-on-Read





The database ruins all our fun

Unless we look at a more document-oriented store for our application

Pay close attention to the database for your deployments

Some of my other courses on this topic:

<https://app.pluralsight.com/library/courses/deploying-databases-octopus>

<https://app.pluralsight.com/library/courses/microsoft-azure-web-applications-services-deploying>



How Many Nines?

Percentage Uptime	Downtime Per Year
99%	Four days
99.9%	Nine hours
99.99%	One hour
99.999%	Five minutes



Our Deployment Strategies

- ~~1. Reconcile the existing state with the changes~~
- ~~2. Blow away what's there and replace it with the changes~~
3. Create a brand-new environment and replace the old one



Two Patterns of Deployment



Snowflake server – a server with a unique configuration



Deploying means reconciling the old state to the new



Immutable Server



Every deployment is the creation, ex nihilo, of an entirely new server

A server whose state cannot change

A server whose state is unrelated to a previous server

Because of that, we can unit, integrate, load, and security test it at our leisure

“With no human intervention”

There was MASSIVE human intervention, just earlier in the process

